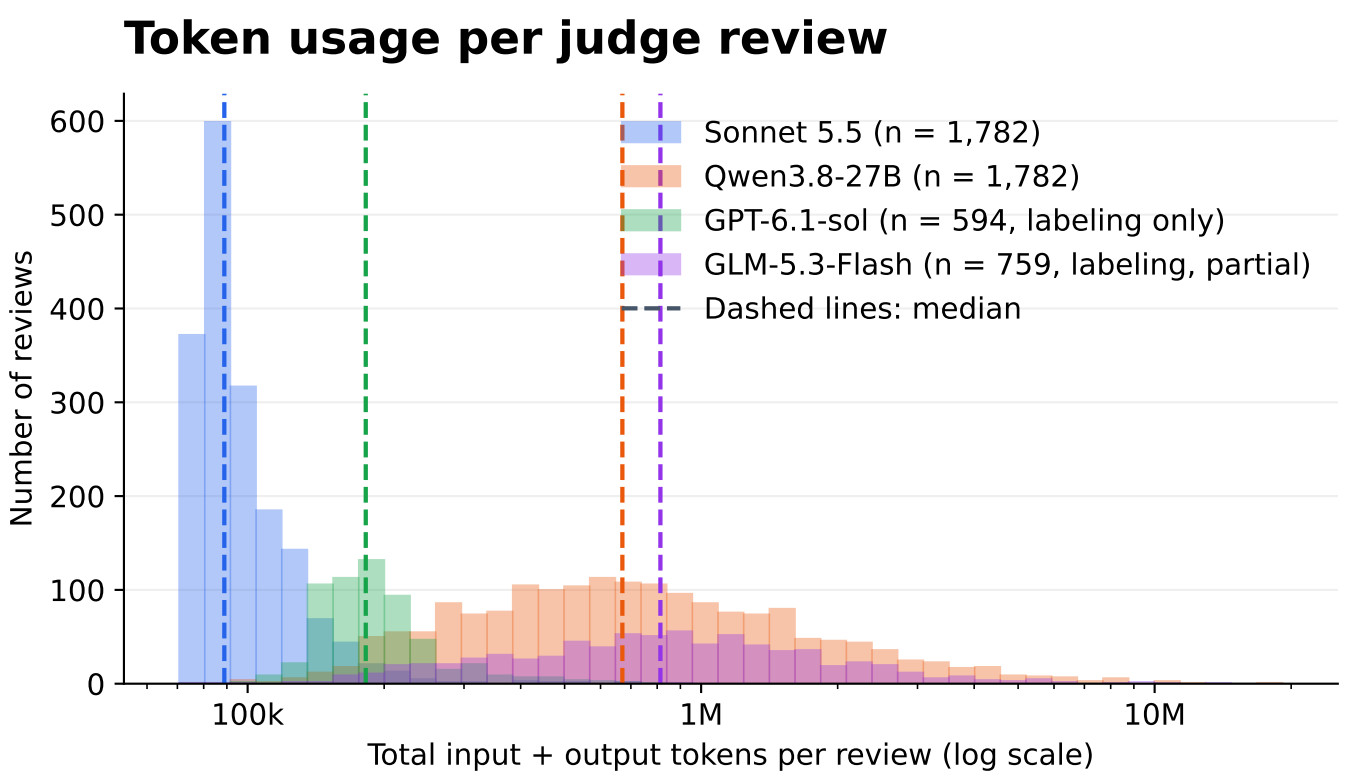
Token usage per judge review
Sonnet 5.5 (n = 1,782)
Qwen3.8-27B (n = 1,782)
GPT-6.1-sol (n = 594, labeling only)
GLM-5.3-Flash (n = 759, labeling, partial)
Dashed lines: median
Number of reviews
0
100
200
300
400
500
600
100k
1M
10M
Total input + output tokens per review (log scale)